

Burrows-Wheeler compression with modified sort orders and exceptions to the MTF phase, and their impact on the compression rate

Marc Lehmann

Burrows-Wheeler compression

- ▶ M. Burrows and D.J. Wheeler in 1994
- ▶ lossless compression
- ▶ context-based
- ▶ 3 stages in basic form

Burrows-Wheeler compression

- ▶ M. Burrows and D.J. Wheeler in 1994
- ▶ lossless compression
- ▶ context-based
- ▶ 3 stages in basic form

Definitions:

symbol smallest logical unit of information

string sequence of symbols

The Burrows-Wheeler Transform

Example: mississippi

- Write all cyclic shifts into a table.

0	m	i	s	s	i	s	s	i	p	p	i
1	i	s	s	i	s	s	i	p	p	i	m
2	s	s	i	s	s	i	p	p	i	m	i
3	s	i	s	s	i	p	p	i	m	i	s
4	i	s	s	i	p	p	i	m	i	s	s
5	s	s	i	p	p	i	m	i	s	s	i
6	s	i	p	p	i	m	i	s	s	i	s
7	i	p	p	i	m	i	s	s	i	s	s
8	p	p	i	m	i	s	s	i	s	s	i
9	p	i	m	i	s	s	i	s	s	i	p
10	i	m	i	s	s	i	s	s	i	p	p

The BWT

- ▶ Sort the table lexicographically to get the *BW table*.
- ▶ Last column is the output of the transform (*BW code*).

0	i	m	i	s	s	i	s	s	i	p	p
1	i	p	p	i	m	i	s	s	i	s	s
2	i	s	s	i	p	p	i	m	i	s	s
3	i	s	s	i	s	s	i	p	p	i	m
4	m	i	s	s	i	s	s	i	p	p	i
5	p	i	m	i	s	s	i	s	s	i	p
6	p	p	i	m	i	s	s	i	s	s	i
7	s	i	p	p	i	m	i	s	s	i	s
8	s	i	s	s	i	p	p	i	m	i	s
9	s	s	i	p	p	i	m	i	s	s	i
10	s	s	i	s	s	i	p	p	i	m	i

The BWT

Context Blocks

- *Context block*: block of BW code corresponding to rows with a specific symbol at the beginning.

0	i	m	i	s	s	i	s	s	i	p	p
1	i	p	p	i	m	i	s	s	i	s	s
2	i	s	s	i	p	p	i	m	i	s	s
3	i	s	s	i	s	s	i	p	p	i	m
4	m	i	s	s	i	s	s	i	p	p	i
5	p	i	m	i	s	s	i	s	s	i	p
6	p	p	i	m	i	s	s	i	s	s	i
7	s	i	p	p	i	m	i	s	s	i	s
8	s	i	s	s	i	p	p	i	m	i	s
9	s	s	i	p	p	i	m	i	s	s	i
10	s	s	i	s	s	i	p	p	i	m	i

The BWT

Reversibility

- ▶ Every column is a permutation of the input
- ▶ First column can be reconstructed by sorting the last

0	i	m	i	s	s	i	s	s	i	p	p	
1	i	p	p	i	m	i	s	s	i	s	s	
2	i	s	s	i	p	p	i	m	i	s	s	
3	i	s	s	i	s	s	i	p	p	i	m	
4	m	i	s	s	i	s	s	i	p	p	i	0
5	p	i	m	i	s	s	i	s	s	i	p	
6	p	p	i	m	i	s	s	i	s	s	i	1
7	s	i	p	p	i	m	i	s	s	i	s	
8	s	i	s	s	i	p	p	i	m	i	s	
9	s	s	i	p	p	i	m	i	s	s	i	2
10	s	s	i	s	s	i	p	p	i	m	i	3

The BWT

Reversibility

- ▶ Start index is also needed
- ▶ i -th occurrence in the first column corresponds to i -th occurrence of the last column

0	i	m	i	s	s	i	s	s	i	p	p	
1	i	p	p	i	m	i	s	s	i	s	s	
2	i	s	s	i	p	p	i	m	i	s	s	
3	i	s	s	i	s	s	i	p	p	i	m	
4	m	i	s	s	i	s	s	i	p	p	i	0
5	p	i	m	i	s	s	i	s	s	i	p	
6	p	p	i	m	i	s	s	i	s	s	i	1
7	s	i	p	p	i	m	i	s	s	i	s	
8	s	i	s	s	i	p	p	i	m	i	s	
9	s	s	i	p	p	i	m	i	s	s	i	2
10	s	s	i	s	s	i	p	p	i	m	i	3

The BWT

The Effect

- ▶ Substrings of the input beginning with the same symbols are sorted one below the other
- ▶ BW code are the symbols immediately preceding them
- ▶ Usually only a few distinct symbols in a context block, many runs of the same symbol

The BWT

The Effect

- ▶ Substrings of the input beginning with the same symbols are sorted one below the other
- ▶ BW code are the symbols immediately preceding them
- ▶ Usually only a few distinct symbols in a context block, many runs of the same symbol

For example, context block corresponding to “nd ” in book1 (first 100 symbols):

```
eaeaaAaaiaaaaaaaaaauaaaaoaaaiaaaiaaauaaauaiaaaaiaaa  
aaaaaaaaaaaaaaaaaaaaiaaaaaAaaaaaaaaaaaaaaaaaaaaa
```

Move-To-Front Coding

- ▶ Alphabet initialized with all possible symbols (e.g. [0x00, 0x01, ..., 0xff] for bytes)
- ▶ Symbols encoded as index in the coder's alphabet
- ▶ Encoded symbols is moved to the front of the alphabet

Move-To-Front Coding

- ▶ Alphabet initialized with all possible symbols (e.g. [0x00, 0x01, ..., 0xff] for bytes)
- ▶ Symbols encoded as index in the coder's alphabet
- ▶ Encoded symbols is moved to the front of the alphabet

For example, encode aaabacccba:

[97, 0, 0, 98, 1, 99, 0, 0, 2, 2]

- ▶ Generally small numbers
- ▶ Runs of same symbols are runs of zeros
- ▶ book1: 0: 49.8%, 1: 15.4%, 2: 8%, 3: 5.3%

Entropy Coding

- ▶ Output of MTF has very skewed probabilities, suitable for entropy coding
- ▶ Symbols with high probability are encoded with short codes
- ▶ Huffman coding, arithmetic coding

Modifying the Sort Order

- ▶ B. Chapin, 1998
- ▶ Instead of $a \rightarrow b \rightarrow c \rightarrow \dots$, sort differently
- ▶ Transitions between “similar” context blocks means lower MTF codes in the beginning
- ▶ Less symbols have to be “fetched from the back” of the alphabet
- ▶ Chapin: handpicked order aeioubcd...
- ▶ Overhead: $\lceil \log_2 256! \rceil = 1684$ bits

Computing Orders

- ▶ Assign a cost to each transitions between symbols (i.e. What if x was sorted before y ?)
- ▶ Run Traveling Salesman Heuristic on the costs
- ▶ Best tour is the best sort order
- ▶ (according to the metric that computed the costs)

Metrics

- ▶ Chapin: based on BW code
- ▶ Analyze similarities in symbol frequencies of context blocks
- ▶ Badness metric: based on the effect on the MTF code
- ▶ Attempts to put a number to how “bad” a transition is (for compression)

The Badness Metric

Partial MTF

- ▶ Like regular MTF, but start with empty alphabet
- ▶ Encode symbols that aren't in it with a special code (-1)

Example aaabacccba from earlier:

MTF: [97, 0, 0, 98, 1, 99, 0, 0, 2, 2]

Partial MTF: [-1, 0, 0, -1, 1, -1, 0, 0, 2, 2]

- ▶ Special codes only difference from regular MTF

The Badness Metric

- ▶ Want to determine badness value for transition from x to y
- ▶ Do BWT with natural order and get the context blocks corresponding to x and y
- ▶ Create the partial MTF for context block y and for the concatenation of both

The Badness Metric

- ▶ Want to determine badness value for transition from x to y
- ▶ Do BWT with natural order and get the context blocks corresponding to x and y
- ▶ Create the partial MTF for context block y and for the concatenation of both

For example, `abc` and `ccaadb`:

right side: [**-1**, 0, **-1**, 0, **-1**, **-1**]

combined: [-1, -1, -1, 0, 0, 2, 0, -1, 3]

- ▶ Metric assumes that context blocks remain unchanged
- ▶ Only positions where the right side has special codes are relevant

The Badness Metric

- Compare values at the relevant positions to ideal values

abc → ccaadb

right side: [-1, 0, -1, 0, -1, -1]

combined: [-1, -1, -1, 0, 0, 2, 0, -1, 3]

ideal: [0, 1, 2, 3]

The Badness Metric

- ▶ Compare values at the relevant positions to ideal values

abc → ccaadb

right side: [-1, 0, -1, 0, -1, -1]

combined: [-1, -1, -1, 0, 0, 2, 0, -1, 3]

ideal: [0, 1, 2, 3]

- ▶ Special code in the combined code: symbol only appears in the right side
- ▶ No information: assume ideal
- ▶ Badness value is the sum of the differences between actual and ideal value

The Badness Metric

Variants

- ▶ Weighting: divide value by the number of special codes in the right side
- ▶ So (long) blocks with many different symbols aren't punished
- ▶ MTF prediction: instead of assuming ideal code, make a guess
 - generic** mean of all MTF codes greater or equal the ideal code
 - specific** mean of all MTF codes greater or equal the ideal code, that are encoding the same underlying symbol

The Badness Metric

First Column Only

- ▶ Metric assumes context blocks remain unchanged, but this is not true
- ▶ When a different order is used, the blocks will also be sorted differently
- ▶ More specific problem: only look for order for first column, rest is ordered with the natural order
- ▶ This way, context blocks stay as they were, metric's assumption holds

Performance of the Metrics

book1

Metric	weighted	MTF prediction	all columns	first column
Badness	✗	✗	-0.04	0.01
	✗	generic	-0.02	0.02
	✗	specific	0.04	0.01
	✓	✗	0.03	0.02
	✓	generic	0.06	0.02
	✓	specific	0.08	0.02
"aeiou. . ."			0.08	0.01
histogram differences			0.04	0.01
number of inversions			0.04	0.01
number of inversions log			0.04	0.01

File: book1, size 6150168 bits.

Performance of the Metrics

paper1

Metric	weighted	MTF prediction	all columns	first column
Badness	✗	✗	-0.40	0.02
	✗	generic	-0.21	0.10
	✗	specific	-0.15	0.10
	✓	✗	-0.24	0.12
	✓	generic	-0.10	0.14
	✓	specific	-0.04	0.13
"aeiou. . ."			0.17	0.02
histogram differences			0.05	0.05
number of inversions			0.11	0.11
number of inversions log			-0.04	0.09

File: paper1, size 425288 bits.

Performance of the Metrics

Observations

- ▶ Not much of a difference
- ▶ More relative improvement with smaller file size (number of transitions)
- ▶ Badness is good for first column, not always for all columns
- ▶ Considering overhead, paper1 actually gets bigger

Multiple Sort Orders

- ▶ Make separate orders for the first two (or more) columns
- ▶ One order: for every distinct symbol in the first column, there's one context block
- ▶ Two orders: for every distinct symbol in the first column, for every distinct symbol in the second that follows it, there's one context block
- ▶ More transitions to optimize, hopefully more compression gains

Multiple Sort Orders

The BWT

0	i	s	s	i	s	s	i	p	p	i	m
1	i	s	s	i	p	p	i	m	i	s	s
2	i	p	p	i	m	i	s	s	i	s	s
3	i	m	i	s	s	i	s	s	i	p	p
4	m	i	s	s	i	s	s	i	p	p	i
5	p	p	i	m	i	s	s	i	s	s	i
6	p	i	m	i	s	s	i	s	s	i	p
7	s	s	i	s	s	i	p	p	i	m	i
8	s	s	i	p	p	i	m	i	s	s	i
9	s	i	s	s	i	p	p	i	m	i	s
10	s	i	p	p	i	m	i	s	s	i	s

- In general, different orders for all different symbols in the first column

Reversibility

- ▶ Can't show reversibility with arbitrary number of orders
- ▶ Can show for two orders
- ▶ Problem: i -th occurrence in first column doesn't correspond to i -th occurrence in last column anymore
- ▶ Solution: "Look ahead" and reorder according to the second column

Reversibility

Looking Ahead

0	i	s	s	i	s	s	i	p	p	i	m	
1	i	s	s	i	p	p	i	m	i	s	s	
2	i	p	p	i	m	i	s	s	i	s	s	
3	i	m	i	s	s	i	s	s	i	p	p	
4	m	i	s	s	i	s	s	i	p	p	i	3
5	p	p	i	m	i	s	s	i	s	s	i	2
6	p	i	m	i	s	s	i	s	s	i	p	
7	s	s	i	s	s	i	p	p	i	m	i	0
8	s	s	i	p	p	i	m	i	s	s	i	1
9	s	i	s	s	i	p	p	i	m	i	s	
10	s	i	p	p	i	m	i	s	s	i	s	

- Get possible indices and sort based on following symbols according to the second order

Performance

Metric	weighted	MTF prediction	all columns	first columns
Badness	✗	✗	-0.06	0.07
	✗	generic	-0.02	0.09
	✗	specific	-0.03	0.09
	✓	✗	-0.06	0.09
	✓	generic	-0.03	0.10
	✓	specific	-0.04	0.09
"aeiou. . ."			0.08	0.01
histogram differences			-0.04	0.03
number of inversions			0.02	0.06
number of inversions log			0.01	0.06

File: book1, size 6150168 bits.

Performance

Observations

- ▶ Better compression than one order if only the first columns are reordered
- ▶ Orders for the second column even less suitable as default order
- ▶ Overhead for storing the orders ($83 \cdot 1684$ bits) outweighs compression gain
- ▶ First columns only requires three orders (natural order as default), not sure if reversible

Exceptions to the MTF

- ▶ Context block “nd ” in book1:

```
eaeaaAaaiaaaaaaaaaauaaaaoaaaaiaaaiaauaauaiaaaaiauaa  
aaaaaaaaaaaaaaaaaaaaiaaaaaAaaaaaaaaaaaaaaaaaaaaa
```

- ▶ Context block “. ”:

```
leyyytetylnyyykrnytehnnyadyyyrkesyyydalsyyyddlednyyyd  
trgkryesendydnekyayswnregsyrmdeycs.nntyhkdegyd!
```

- ▶ Next sentence doesn't give indication about last letter of last word of previous one
- ▶ Many different symbols, no long runs

Exceptions to the MTF

- ▶ Symbols are the last letters of words, different probabilities
- ▶ But no information is taken from the further context
- ▶ Encoding with MTF doesn't make sense in this case
- ▶ “Pollutes” the statistics of a static entropy coder

Exceptions to the MTF

- ▶ Symbols are the last letters of words, different probabilities
- ▶ But no information is taken from the further context
- ▶ Encoding with MTF doesn't make sense in this case
- ▶ “Pollutes” the statistics of a static entropy coder
- ▶ Exclude certain blocks from the MTF phase
- ▶ Put a special code in the MTF to signalize a missing block
- ▶ Encode with Huffman directly and append

Exceptions to the MTF

Selecting Exceptions

- ▶ Encode with MTF as usual
- ▶ For every context block: if the mean of all the MTF values is above a threshold, exclude it from the MTF phase
- ▶ Also require a certain length of the block, so the compression gains aren't destroyed by the overhead

Exceptions to the MTF

Performance, book1

min length	threshold	gain	excepted blocks
0	3	-0.39	<i>many</i>
0	4	0.59	0x00, \n, 0x1a, space, !, &,), *, +, ",", ., 0, 5, 7, :, ;, =, >, ?, E, U, V, X
100	4	0.59	\n, space, !, +, ",", ., :, ;, >, ?, E, U
100	4.5	0.60	\n, space, !, +, ",", ., :, ;, >, ?
100	5	0.59	\n, space, !, +, ",", ., :, ;, >
100	6	0.26	\n, ",", .

File: book1, size 6150168 bits.

Exceptions to the MTF

Observations

- ▶ Works much better than the reordering stuff
- ▶ Good choice for threshold seems to be between 4 and 5

Exceptions to the MTF

Observations

- ▶ Works much better than the reordering stuff
- ▶ Good choice for threshold seems to be between 4 and 5

Caveats

- ▶ Effective with static (two-pass) Huffman coder I use
- ▶ Adaptive coder could adapt to temporarily higher MTF codes
- ▶ Huffman coder with multiple tables could have one for these bad cases